




# Linux Day Milano 2022

# Ansible

---

The recipe for a good server



POLITECNICO OPEN  
unix LABS

**Luca Biscaldi**

*bisca@poul.org*

# What is Ansible?



# What is Ansible?

**Ansible** is a *free and open-source* software that **automates**:

- **Provisioning**: preparing an IT infrastructure to execute specific tasks.
- **Configuration management**: maintaining the infrastructure in the desired operating state, even after updates or changes.
- **Applications deploy**: installation, configuration and startup of applications.
- **Orchestration**: coordination between IT systems.

Ok, all nice, but...

**It's not like I have to manage a  
datacenter!**

What can I do with Ansible?

# When can Ansible be useful *to me*?

Some realistic use cases:

- I have to **configure from scratch** a new machine, but it's the tenth time I install Linux and I've had enough.
- I have a machine that I have been using for a long time and I want to **do a clean reinstall**, but I don't remember how I configured it two years ago.

## When can Ansible be useful to me

- I have to **configure a lot of containers** in my new server and I don't want to do it manually.
- I've already configured my machine but I want it to remain in a known **steady state over time**.

Why can't I use a *simple* script on every machine?

## Ansible vs script

1. A script executes **actions** on single machines,  
**Ansible** makes sure that the entire infrastructure is in a  
specific **state**.



## Example 1: Ansible vs script

*I want the file A to contain the line B*

**Script:**

Check that the file exists (otherwise create it),  
Check that it contains the line (otherwise insert it).

**Ansible:** File A *must* contain line B.

2. **Ansible** uses a more **readable** language than a script.

## Example 2: Ansible vs script

*I want the file A to contain the line B.*

### Script:

```
#!/bin/bash
FILE_PATH=A
LINE=B
if [ ! -f "$FILE_PATH" ]; then
    touch $FILE_PATH
fi
if ! grep -Fxq "$LINE" $FILE_PATH; then
    echo "$LINE" >> $FILE_PATH
fi
```

### Ansible:

```
- lineinfile:
  path: A
  line: B
  create: yes
```

## Ansible vs script

3. Ansible allows the parallel management of more machines.

## Example 3:      Ansible vs script

*I want both A and B packages to be installed  
in HOST\_A and HOST\_B machines,  
but only if the Linux distribution is Ubuntu.*

### Ansible:

```
- hosts: HOST_A, HOST_B
  tasks:
    - apt:
      pkg:
        - A
        - B
      when: ansible_distribution == 'Ubuntu'
```

### Script:

```
askStackOverflow();
```

## Ansible strenghts:

- **Minimal**: It does not need the target machines to have specific software installed (only python and a SSH server\*).
- **Secure** : All\* operations are transmitted using SSH.

- **Simple**: The files that make up our project are **straightforward** to read.
- **Idempotent**: Successfully executed operations will always have the *same result*, regardless of how many times they are executed, or the initial state of the machine.
- **Scalable**: The number of managed machines does not influence the project maintenance difficulty.

I'm convinced,

**Let's see how it works!**



## Before starting...

**Ansible** has to be installed only on one machine, called **control node**, that has to be able to communicate with all the machines we want to manage, called **controlled nodes**.

On the controlled nodes it is sufficient that python and a SSH server\* are installed, nothing more is needed!

## ...Let's install it!

In order to **install Ansible** we can:

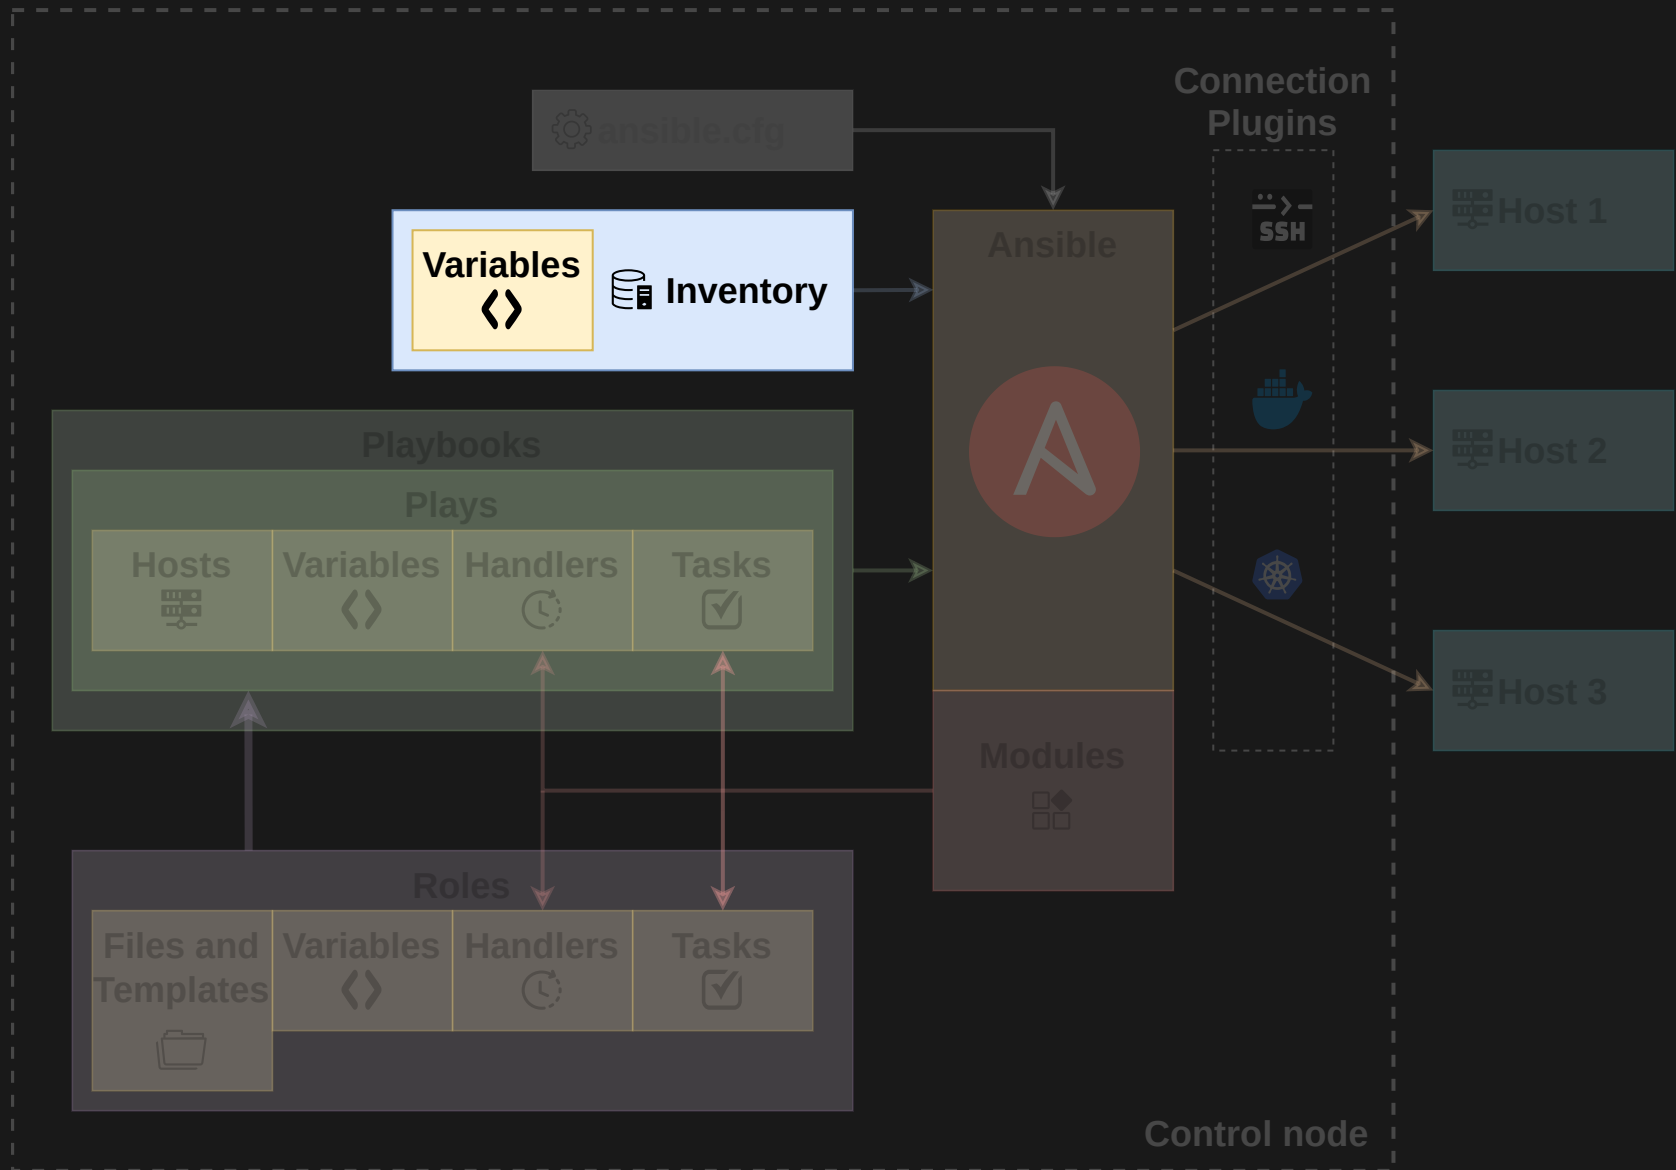
- Use the **package manager** of our GNU/Linux distribution.
- Use **pip**:

```
$ pip install ansible
```

- Compile it from **source**.

How do we "introduce" our other machines to **Ansible**?

# Operating architecture



# Inventory

The **inventory** is the component in our project that defines the infrastructure Ansible will operate on.

It can be a single file (**static inventory**) with simple infrastructures, but we can also use more files, also editable by external services (**dynamic inventory**).

The default inventory is `/etc/ansible/hosts` (single file).

# Content

## Inventory

Inside the inventory we can define:

- **Hosts**: target machines on which **Ansible** will act.
- **Groups**: logical "containers" to group hosts (or othergroups).
- **Variables**: properties associated to groups or hosts.

# Format

# Inventory

An inventory can be written in **INI** syntax or in **YAML** syntax.



**INI** format is linear and more convenient for simple **Inventory** infrastructures:

```
mailserver.example.com db_backend=db_three.example.com
```

```
[webservers]
```

```
foo.example.com db_backend=db_one.example.com
```

```
bar.example.com db_backend=db_two.example.com
```

```
[webservers:vars]
```

```
http_port=80
```

```
[dbservers]
```

```
db_one.example.com
```

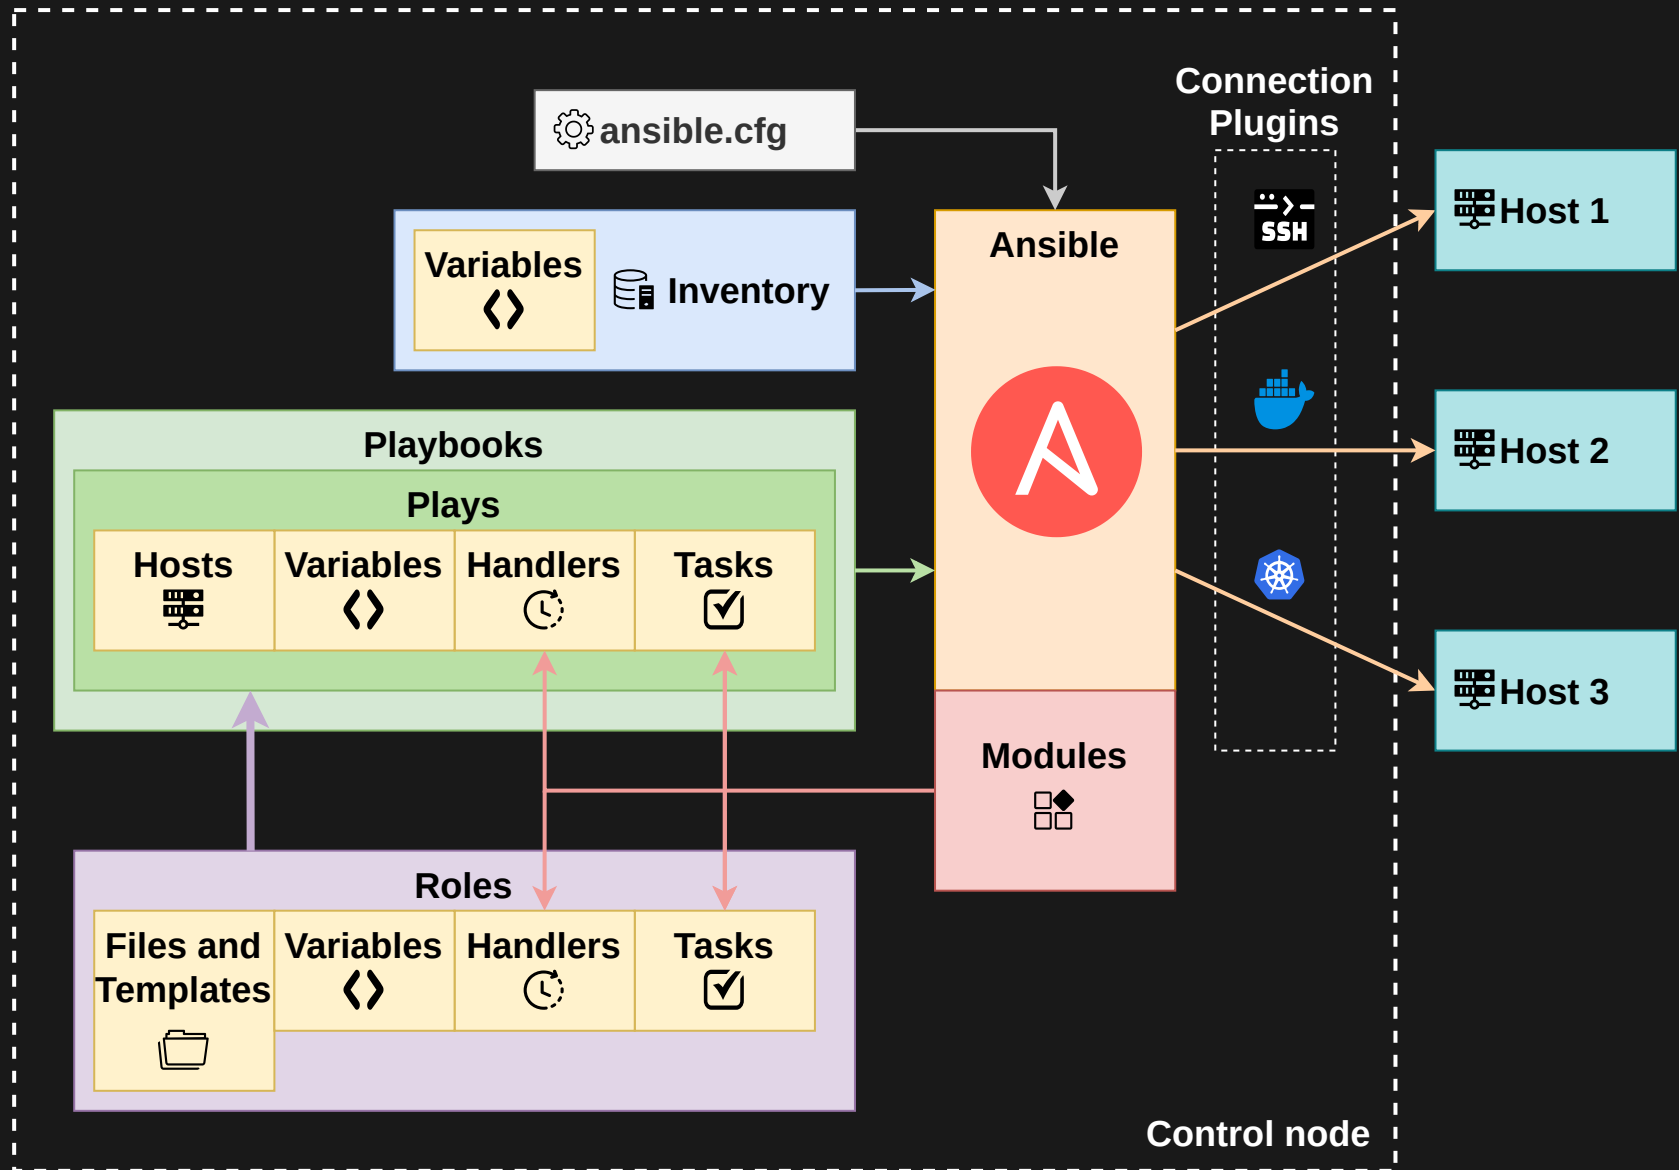
```
db_two.example.com
```

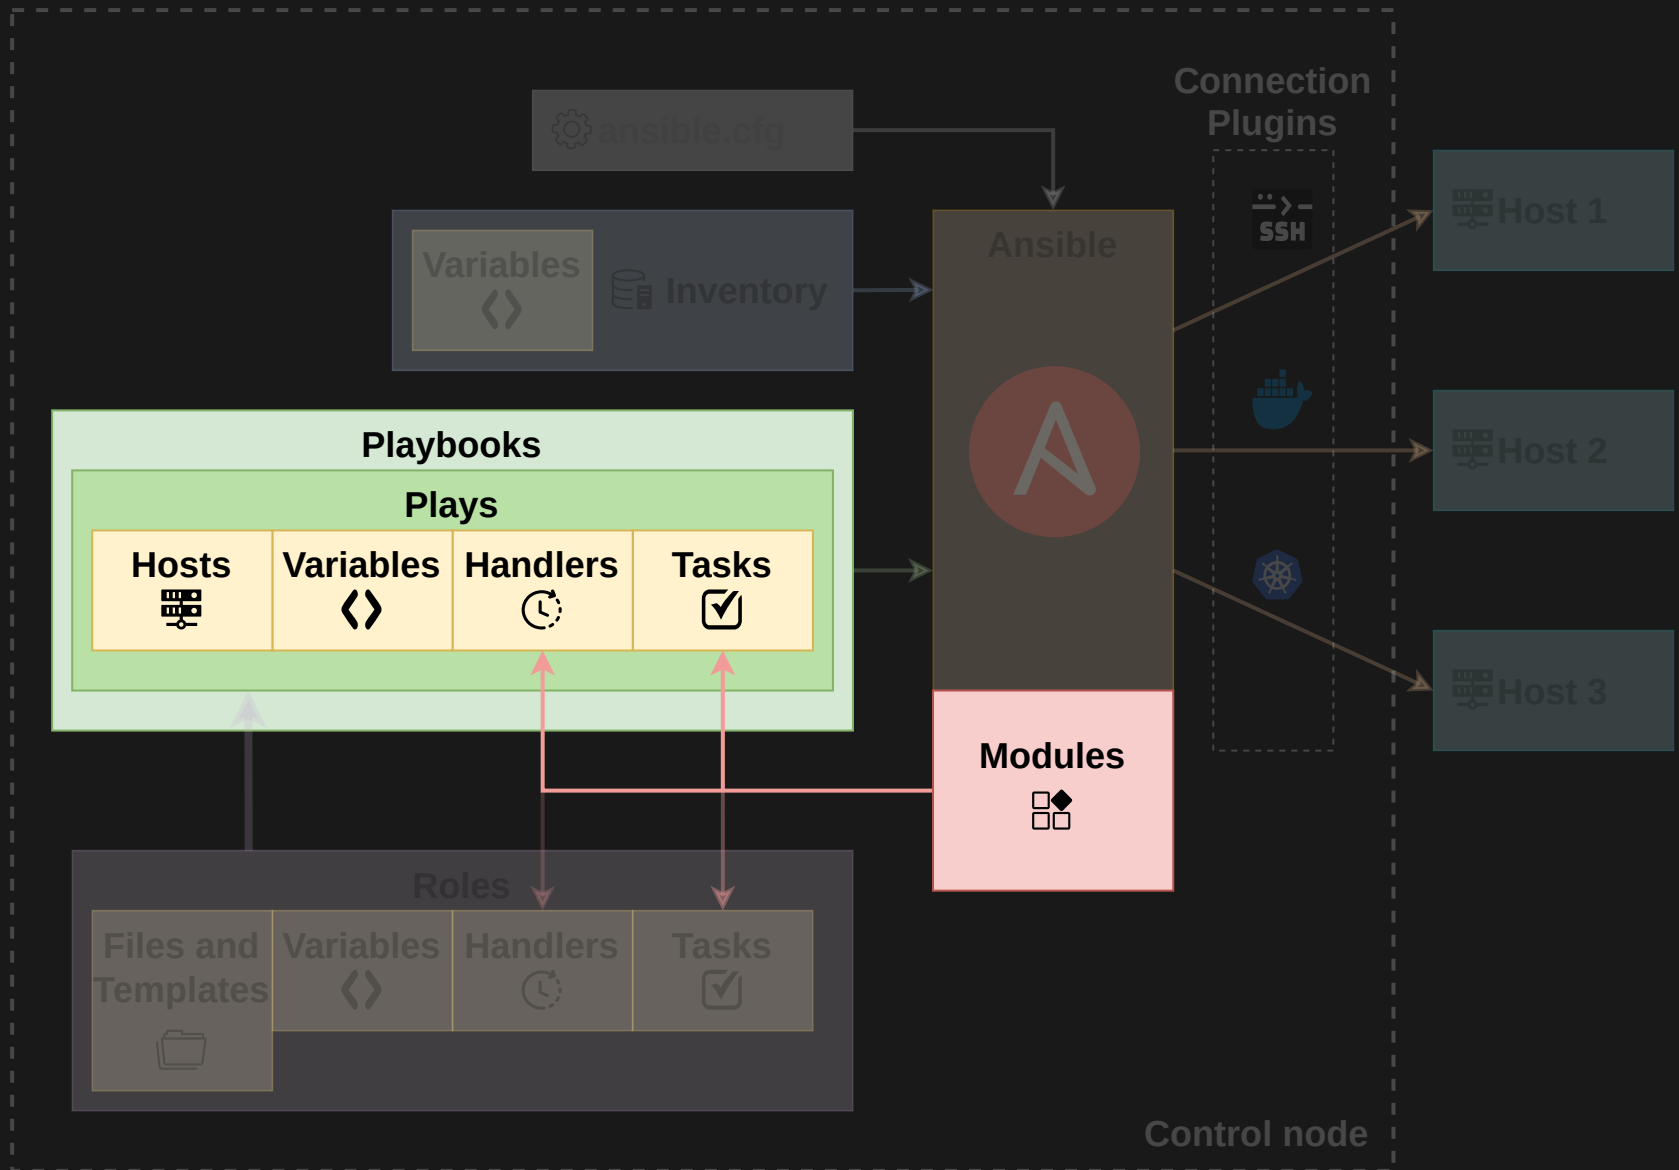
```
db_three.example.com
```

```
[ubuntu:children]
```

```
webservers
```

**Let's do something useful!**





# Modules

**Modules** are code units that Ansible executes on remote hosts.

Each module serves a specific function and might require specific parameters in order to be invoked.

Modules can be individually executed: **Modules**

```
$ ansible <hosts> -m module_name -a module_parameters
```

# Task

A **task** is an action unit. It includes (at least):

- The **module** to invoke.
- The **parameters** needed by the module.

# Play

A **play** is a unit that contains all the information needed to execute a task, such as (at least):

- Information on the **hosts** on which we execute the tasks
- Information on the **tasks** to execute



# Playbook

A **playbook** is a file that contains one or more **plays** to run,  
in **YAML** syntax.

# Playbook example

```
---  
- name: Update web servers  
  hosts: webservers  
  
  tasks:  
    - name: Apache is at the latest version  
      ansible.builtin.yum:  
        name: httpd  
        state: latest  
    - name: Apache service is enabled and running  
      ansible.builtin.service:  
        name: httpd  
        state: started  
        enabled: yes
```

A playbook can be run with:

```
$ ansible-playbook <hosts> <my_playbook.yml>
```

# Handlers

A **handler** is a task that will be executed *only if* another task changes something on the remote host.

To call a handler we use the keyword `notify`.

# Example

## Handlers

*Restart Apache daemon only if another task changes its configuration file.*

```
- name: My play
  hosts: webservers

  tasks:
    - name: Apache config file is updated
      ansible.builtin.copy:
        src: /srv/httpd.j2
        dest: /etc/httpd.conf
      notify:
        - Restart apache

  handlers:
    - name: Restart apache
      ansible.builtin.service:
        name: httpd
        state: restarted
```

# Variables

Variables are objects containing information.  
They can be defined:

## Variables

### 1. Inside the **inventory**:

```
all:  
hosts:  
    mailserver.example.com  
vars:  
    my_variable: antani
```

or inside the `host_vars/` and `group_vars/` subfolders.

## Variables

### 2. Inside a **play**:

```
- name: My play
  hosts: webserver
  vars:
    http_port: 80
  tasks:
    ...
```



3. Inside a **role**, in the `vars/` and `defaults/` subfolders. **Variables**

4. From command line, with `-e variable=value` flag. **Variables**

# Using variables

## Variables

Variables can be used by calling them with:

```
"{{ my_variable }}"
```



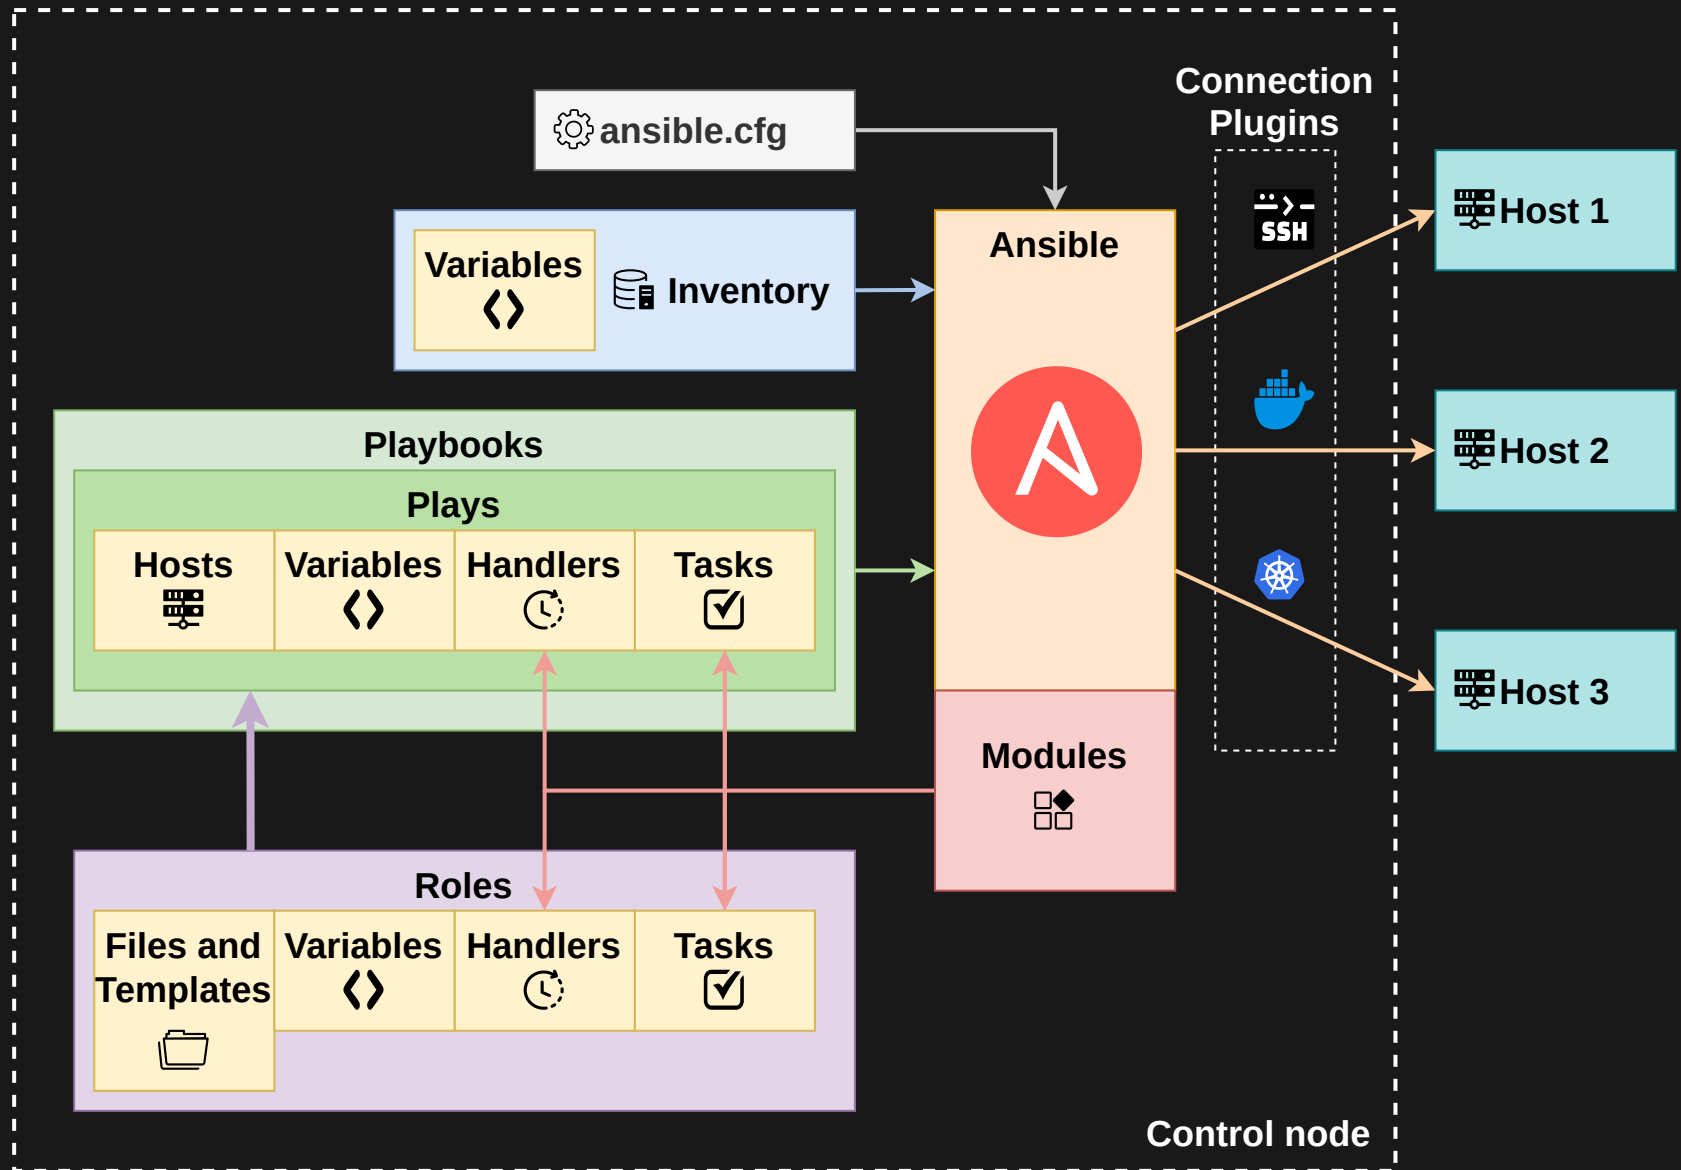
Double quotes are mandatory!

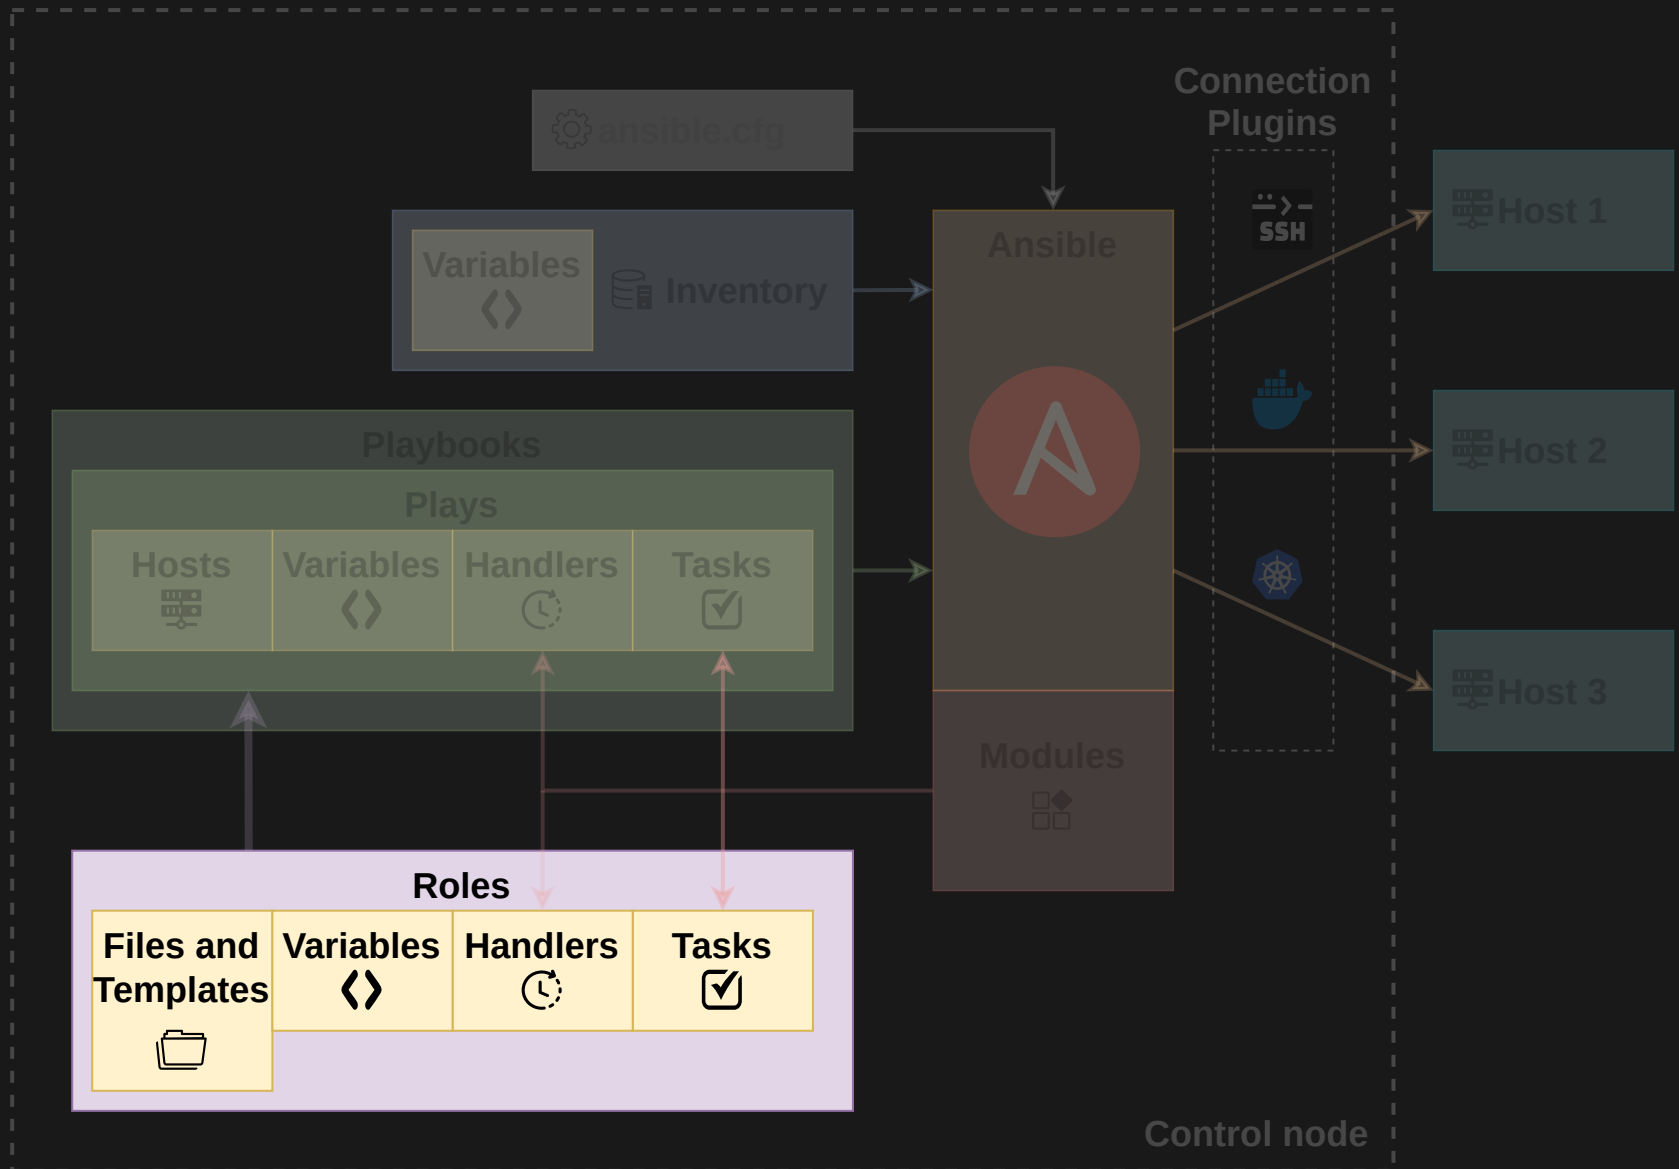
# What if I declare the same variable in different places?

Precedence rules

(last one has precedence over the others):

1. Role defaults
2. Inventory
3. Play
4. Role variables
5. Task
6. Command Line





# Roles

A **role** is a set of:

- Tasks
- Handlers
- Variables
- Files and templates

that can be loaded in a **play**.

# What are roles for?

Roles

Roles are useful to **organize** and to make the project **modular**.



## Roles

Example:

- I write a **common** role with the tasks I want to run on allmachines (i.e. import SSH keys, create users...).
- I write a **webservers** role with the tasks needed for the configuration of a webserver (i.e. install needed packages).

## Roles

My playbook will be:

```
- name: Setup webserver
  hosts: webserver
  roles:
    - common
    - webserver

- name: Setup other machines
  hosts: others
  roles:
    - common
```

## Roles

To use a role we can:

1. Add it in a play in the roles **section**:

```
- hosts: all
  roles:
  - MY_ROLE
```

### 2. **Include** it as a task in a play:

```
- hosts: all
  tasks:
    - name: My beautiful role
      include_role:
        name: MY_ROLE
```

### 3. **Import** it as a task in a play:

```
- hosts: all
  tasks:
    - name: My beautiful role
      import_role:
        name: MY_ROLE
```

# Useful modules

## Setup

## Useful modules

Gather information about remote hosts.

```
$ ansible <hosts> -m setup
```

**Useful modules**  
*Setup* module registers a special variable,  
`ansible_facts`, for each host. It is a dictionary  
containing *a lot of* useful information about the remote host:

```
{
  "ansible_all_ipv4_addresses": [
    "REDACTED IP ADDRESS"
  ],
  "ansible_all_ipv6_addresses": [
    "REDACTED IPV6 ADDRESS"
  ],
  "ansible_apparmor": {
    "status": "disabled"
  },
  "ansible_architecture": "x86_64",
  "ansible_bios_date": "11/28/2013",
  "ansible_bios_version": "4.1.5",
  "ansible_cmdline": {
    "BOOT_IMAGE": "/boot/vmlinuz-3.10.0-862.14.4.el7.x86_64",
    "console": "ttyS0,115200",
    .....
```



Setup module will be run *by default* at the start of every **Useful modules** playbook.

To override this behaviour we can use

```
gather_facts: no
```

inside our play.

## Template

## Useful modules

Create a file based on a template. It's useful when the content of the file I want to create depends on variables.

## Useful modules

Templates (by default) are located in the `[role]/templates/` subfolder and are written in *Jinja2* syntax.

Template example (for `/etc/network/interfaces`):

```
auto {{ interface.name }}
iface {{ interface.name }} inet static
    address {{ interface.address }}/{{ interface.netmask }}
    gateway {{ interface.gateway }}
```

## Useful modules

The task that allows to use this template would be:

```
- name: update network interfaces
  template:
    src: my_template.jn2
    dest: /etc/network/interfaces
```

## Other useful modules

- **copy**: it copies a file from the control node to the controlled node.
- **lineinfile**: it checks that a specified line is present inside a file.
- **file**: it manages files, directories and their properties.

What if I don't want to use SSH as root user?

## Privilege escalation

# Privilege escalation

Changing user, gaining administrative privileges.

# Privilege escalation

## How?

Using 2 parameters in the play:

- Declare that we intend to change user:

```
become: yes
```

- Declare which user we want to be (default is root):

```
become_user: root
```



What did we *not* mention?

I have a lot of secrets, do I keep them under the mattress?

# Ansible Vault

# Ansible Vault

It allows to save sensible information inside our project.

All information will be **encrypted** and a password will be asked every time we run our playbook.

I am *very lazy* and i don't want to write all my playbooks  
from scratch.

## Ansible Galaxy

# Ansible Galaxy

It's an hub where roles and entire projects are shared, from which we can freely draw.

# Ansible Galaxy

## How to use it?

From command line!

```
$ ansible-galaxy <action> <name>
```

- **search**: to search a term.
- **info**: to gather information about a role or a project we found.
- **install**: to install the role or the project locally.

I am *even lazier* and I am tired of using the terminal.

**Ansible AWX**

## Ansible AWX

It's a GUI in the form of a **hostable service** that allows us to manage our Ansible projects.



# Ansible AWX


## Capabilities:

- **Realtime jobs update:** We can have the output of each executed task, for each machine, in real time.
- **Multiple users:** we can have different users, with different permissions, organized in groups.
- **Logs:** we can see who did what.
- **Scheduling:** we can automate our automation.


# Links

- [Ansible Webpage](#)
- [Ansible Docs Page](#)

# Thank you for your attention!



POLITECNICO OPEN  
unix LABS



This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](#).

Source code available [here](#)

**Luca Biscaldi**

*bisca@poul.org*